

Building Common Lisp programs using Bazel

or Correct, Fast, Deterministic Builds for Lisp

James Y. Knight
Google
jyknight@google.com

François-René Rideau
Google
tunes@google.com

Andrzej Walczak
Google
czak@google.com

ABSTRACT

We will demonstrate how to build Common Lisp programs using Bazel, Google's hermetic and reproducible build system. Unlike the state of the art so far for building Lisp programs, Bazel ensures that incremental builds are always both fast and correct. With Bazel, one can statically link C libraries into the SBCL runtime, making the executable file self-contained.

Categories and Subject Descriptors

D.2.3 [Software Engineering]: Coding Tools and Techniques

Keywords

Bazel, Build System, Common Lisp, Determinism, Hermeticity, Reproducibility

1. INTRODUCTION

Common Lisp, a general-purpose programming language, is used at Google for the well-known low-fare search engine QPX (de Marcken 2001). Google now builds its Lisp code incrementally, using Bazel, its recently open-sourced scalable build system.

Bazel is designed to build software in a reproducible and hermetic way. Hermeticity means all build dependencies, including build tools such as compilers, are kept under source control. Bazel can thus assess what was or wasn't modified, and either rebuild or reuse cached build artifacts. Reproducibility means building the same target multiple times from the same source code produces the same output. Reproducibility facilitates testing and debugging production code. Bazel further enforces determinism by executing each build action in a container wherein only declared inputs may be read, and any non-declared output is discarded. Bazel can thus parallelize build actions, either locally or to remote workers. Bazel supports writing software in a mix of languages notably including C++, Java, Python, Go and Javascript. Compilers and other build tools must be tuned to remove sources of non-determinism such as timestamps, PRNG seeds, etc.

While mainly written in Java, Bazel is extensible using *Skylark*, a subset of Python with strict limits on side-effects. We used Skylark

to add support for building software written in Common Lisp.

2. PREVIOUS WORK

The state of the art so far for building large Common Lisp applications is ASDF (Rideau 2014). A descendent of the original Lisp DEFSYSTEM from the 1970s, ASDF builds all the code in the current Lisp image; incremental builds may therefore be affected by all kinds of potential side-effects in the current image; and to guarantee a deterministic build one has to build from scratch. ASDF also lacks good support for multi-language software. An attempt to build Lisp code deterministically, XCVB (Brody 2009), failed for social and technical reasons, though it had a working prototype.

Meanwhile, QPX was built using an ad-hoc script loading hundreds of source files before compiling them and reloading the resulting FASLs. The multi-stage build was necessary because of circular dependencies between the files; the dependencies formed a big "hairball" which any replacement build solution had to handle.

3. BUILDING LISP CODE WITH BAZEL

The input to Bazel is a set of BUILD files with Python-like syntax, that declaratively specify software *targets* using *rules*.

Our first Lisp support function, `lisp_library`, declares an intermediate target which can be referenced from other Lisp rules. With SBCL (Steel Bank Common Lisp), `lisp_library` creates a FASL Load (FASL) archive by concatenating the FASL files produced by compiling each of its Lisp sources.

Bazel users specify *attributes* when defining rules. For Lisp rules, these include the Lisp sources `srcs`, Lisp libraries `deps`, C sources `csrcs` and C libraries `cdeps`, auxiliary data available at runtime to all executable targets depending on the library, and auxiliary `compile_data` available at compile-time.

Lisp rules have additional build options. The `order` attribute notably specifies a build strategy; the default "serial" order loads each source file in sequence before compiling the next file. The "parallel" order compiles all files in parallel without loading other ones. The "multipass" order first loads all source files, then compiles each one separately in parallel, which is useful to handle a "hairball" aggregate.

```
load("@lisp_bazel//:bazel/rules.bzl",
     "lisp_library")
lisp_library(
    name = "alexandria",
    srcs = ["package.lisp",
           # ...
           "io.lisp"],
    visibility = ["//visibility:public"])
```

The above example is from the BUILD file for the "alexandria" general utility library. Bazel first loads `lisp_library` from

its conventional *build label* under the `lisp__bazel` "external repository". The `visibility` attribute indicates which Bazel packages are allowed to reference the rule's target — in this case, all packages.

The following command builds `alexandria.fasl` and makes it available at a well defined path:

```
bazel build :alexandria
```

Our second function, `lisp_binary`, creates an executable including both Lisp runtime and Lisp core image. If Lisp or C sources are specified, they will be compiled into corresponding Lisp and C components before being statically linked into the final binary. Our third function, `lisp_test`, is a variation on the `lisp_binary` rule meant to be invoked with the `bazel test` command.

```
load("@lisp__bazel//:bazel/rules.bzl",
     "lisp_binary")
lisp_binary(
  name = "myapp",
  srcs = ["myapp.lisp"],
  main = "myapp:main",
  deps = [
    "@lisp__alexandria//:alexandria"])
```

This BUILD file contains a `lisp_binary` target which references the "alexandria" BUILD target seen before. When running the binary, the Lisp function `myapp:main` will be called with no arguments at startup. The program may be compiled and executed using:

```
bazel run :myapp
```

A `lisp_binary` can directly or transitively depend on C or C++ libraries. Static linking of the libraries makes it more reliable to deploy such a binary on multiple hosts in the cloud, without the opportunity to get library dependencies wrong; in particular, it helps with minimizing discrepancies between test and production environments. C and C++ dependencies can be specified via the `cdeps` rule attribute, which can refer to any `cc_library` built with Bazel. The `csrcs` and `copts` rule attributes allow to directly specify C source files for which an internal target will be generated.

Thanks to these build rules, the duration of the incremental QPX build went from about 15 minutes to about 90 seconds, with qualitative effects on developer experience. However, this is for a large project, using a computing cloud for compilation. The open source version of Bazel currently lacks the ability to distribute builds, but it can already take advantage of multiple cores on a single machine. The typical Lisp user will therefore not experience as large a speedup when using the Bazel lisp rules.

4. INSIDE THE LISP RULES

Lisp support was implemented using Bazel's *Skylark* extension language. The functions `lisp_library`, `lisp_binary` and `lisp_test` are Skylark *macros* calling internal implementation *rules*. A Skylark *macro* is basically a Python function that is executed by Bazel at the time the BUILD file is loaded and invokes the actual Skylark rules as side-effects. A Skylark *rule* consists of an implementation function and a list of attribute specifications that notably define type-checked inputs and outputs for the rule's target. The Lisp support macros establish two separate graphs for each of the Lisp and C parts of the build, that are connected at the final binary targets. Thus, the `lisp_library` macro calls the `_lisp_library` rule to create Lisp related actions, and also calls the `make_cdeps_library` macro to create the C related targets using Skylark's `native.cc_library`.

The rules compile C sources and Lisp sources in parallel to each other, and the resulting compilation outputs are combined together in the last step. This improves the build parallelism and reduces

the latency. In order to facilitate linking, all C symbols referred to at Lisp-compilation time are dumped into a linker script file. The final linking step uses that `.lds` file to include from C libraries only the referenced objects, and to statically detect any missing or misspelled C symbol. The `lisp_binary` and the `lisp_test` macros then create the executable by combining the runtime of SBCL linked with additional C libraries and a core image dumped after loading all FASLS.

The `_lisp_library` rule implementation compiles each file in a new Lisp process (possibly on remote worker machines) that will first load all the Lisp *source* files from all the transitive dependencies as well as relevant files from the current rule (depending on the `order` attribute). This is faster than loading FASLS, thanks to SBCL's *fasteval* interpreter (Katzman 2015), written specifically to speed up building with Bazel; this also enables all compile actions to run in parallel, whereas loading FASLS would introduce long chains of dependencies between actions and cause high latency. On the downside, some source files must be fixed to add missing `:execute` situations in `eval-when` forms, and optionally to explicitly compile computation-intensive functions used at compile-time.

The `_lisp_library` implementation also returns a *provider* structure containing transitive information about: all Lisp sources loaded; all reader features declared; runtime and compile-time data for each library; FASLS from each `lisp_library` target, used when linking the final binary target; deferred warnings from each compilation — mostly for undefined functions — checked after all FASLS have been loaded into the final target.

5. REQUIREMENTS

The Lisp support for Bazel so far only works on SBCL, on the x86-64 architecture, on Linux and MacOS X. It shouldn't be hard to make it work on a platform supported by both SBCL and Bazel. However, porting to another Lisp implementation will be non-trivial, notably with respect to linking C libraries statically or to achieving latency as low as with the *fasteval* interpreter.

Bazel itself is an application written in Java. It takes seconds to start for the first time; then it becomes a server that can start an incremental build instantly but consumes gigabytes of memory.

6. CONCLUSION AND FUTURE WORK

We have demonstrated simultaneously how Common Lisp applications can be built in a fast and robust way, and how Bazel can be extended to reasonably support a new language unforeseen by its authors. Bazel may not be a lightweight solution for writing small programs in Lisp; but it is a proven solution for building large industrial software projects tended by several groups of developers using multiple programming languages, including Lisp.

In the future, we may want to add Lisp-side support for interactively controlling Bazel: we would like to be able to build code, and load the result code into the current image, without reloading unmodified FASLS and object files.

Our code is at: <http://github.com/qitab/bazelisp>

Bibliography

- François-René Rideau and Spencer Brody. XCVB: an eXtensible Component Verifier and Builder for Common Lisp. 2009.
- Carl de Marcken. Carl de Marcken: Inside Orbitz. 2001. <http://www.paulgraham.com/carl.html>
- Douglas Katzman. SBCL's Fasteval interpreter. 2015.
- François-René Rideau. ASDF3, or Why Lisp is Now an Acceptable Scripting Language (extended version). 2014.